

Final Project

This notebook is adapted from here:

https://aiqm.github.io/torchani/examples/nnp_training.html

Checkpoint 1: Data preparation

1. Create a working directory:

`/global/scratch/users/[USER_NAME]/[DIR_NAME]` . Replace the `[USER_NAME]` with yours and specify a `[DIR_NAME]` you like.

2. Copy the Jupyter Notebook to the working directory
3. Download the ANI dataset `ani_dataset_gdb_s01_to_s04.h5` from bCourses and upload it to the working directory

```
In [1]: import warnings
warnings.filterwarnings("ignore", category=UserWarning)
import numpy as np
from tqdm import tqdm
import torch
import torch.nn as nn
import torchani

print(torch.__version__)
print(torchani.__version__)
```

2.2.2

2.2.4

Use GPU

```
In [2]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)
```

cuda

Set up AEV computer

AEV: Atomic Environment Vector (atomic features)

Ref: Chem. Sci., 2017, 8, 3192

```
In [3]: def init_aev_computer():
    Rcr = 5.2
    Rca = 3.5
    EtaR = torch.tensor([16], dtype=torch.float, device=device)
    ShfR = torch.tensor([
        0.900000, 1.168750, 1.437500, 1.706250,
        1.975000, 2.243750, 2.512500, 2.781250,
        3.050000, 3.318750, 3.587500, 3.856250,
        4.125000, 4.393750, 4.662500, 4.931250
    ], dtype=torch.float, device=device)

    EtaA = torch.tensor([8], dtype=torch.float, device=device)
    Zeta = torch.tensor([32], dtype=torch.float, device=device)
    ShfA = torch.tensor([0.90, 1.55, 2.20, 2.85], dtype=torch.float, device=device)
    ShfZ = torch.tensor([
        0.19634954, 0.58904862, 0.9817477, 1.37444680,
        1.76714590, 2.15984490, 2.5525440, 2.94524300
    ], dtype=torch.float, device=device)

    num_species = 4
    aev_computer = torchani.AEVComputer(
        Rcr, Rca, EtaR, ShfR, EtaA, Zeta, ShfA, ShfZ, num_species
    )
    return aev_computer

aev_computer = init_aev_computer()
aev_dim = aev_computer.aev_length
print(aev_dim)
```

384

Prepare dataset & split

```
In [4]: def load_ani_dataset(dspath):
    self_energies = torch.tensor([
        0.500607632585, -37.8302333826,
        -54.5680045287, -75.0362229210
    ], dtype=torch.float, device=device)
    energy_shifter = torchani.utils.EnergyShifter(None)
    species_order = ['H', 'C', 'N', 'O']

    dataset = torchani.data.load(dspath)
    dataset = dataset.subtract_self_energies(energy_shifter, species_order)
    dataset = dataset.species_to_indices(species_order)
    dataset = dataset.shuffle()
    return dataset

dataset = load_ani_dataset("ani_gdb_s01_to_s04.h5")
# Use dataset.split method to do split
train_data, val_data, test_data = dataset.split(.8, .1, .1)
```

Batching

```
In [5]: batch_size = 8192
# use dataset.collate(...).cache() method to do batching
train_data_loader = train_data.collate(batch_size).cache()
val_data_loader = val_data.collate(batch_size).cache()
test_data_loader = test_data.collate(batch_size).cache()
```

Torchani API

```
In [6]: class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 128),
            nn.ReLU(),
            nn.Linear(128, 1)
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [7]: # train_data_batch = next(iter(train_data_loader))

for train_data_batch in train_data_loader:
    loss_func = nn.MSELoss()
    species = train_data_batch['species'].to(device)
    coords = train_data_batch['coordinates'].to(device)
    true_energies = train_data_batch['energies'].to(device).float()
    _, pred_energies = model((species, coords))
    loss = loss_func(true_energies, pred_energies)
    print(loss)
    break
```

```
tensor(0.0355, device='cuda:0', grad_fn=<MseLossBackward0>)
```

Checkpoint 2

```
In [8]: coords.shape[0]
train_epoch_loss = 0.0

loss_func = nn.MSELoss()
batch_loss = loss_func(true_energies, pred_energies)
batch_importance = coords.shape[0] / len(train_data)
train_epoch_loss += batch_loss.detach().cpu().item() * batch_importance

print(train_epoch_loss)

0.0004199363558073355
```

```
In [9]: from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
```

```
In [10]: class ANITrainer:
    def __init__(self, model, batch_size, learning_rate, epoch, l2):
        self.model = model

        num_params = sum(item.numel() for item in model.parameters())
        print(f"{model.__class__.__name__} - Number of parameters: {num_params}")

        self.batch_size = batch_size
        self.optimizer = torch.optim.Adam(model.parameters(), learning_rate,
        self.epoch = epoch

    def train(self, train_data, val_data, early_stop=True, draw_curve=True):
        self.model.train()

        # init data loader
        print("Initialize training data...")
        train_data_loader = train_data.collate(batch_size).cache()

        # definition of loss function: MSE is a good choice!
        loss_func = nn.MSELoss()

        # record epoch losses
        train_loss_list = []
        val_loss_list = []
        lowest_val_loss = np.inf

        for i in tqdm(range(self.epoch), leave=True):
            train_epoch_loss = 0.0
            for train_data_batch in train_data_loader:

                # compute energies
                species = train_data_batch['species'].to(device)
                coords = train_data_batch['coordinates'].to(device)
                true_energies = train_data_batch['energies'].to(device).float
```

```

        _, pred_energies = model((species, coords))

        # compute loss
        batch_loss = loss_func(true_energies, pred_energies)

        # do a step
        self.optimizer.zero_grad()
        batch_loss.backward()
        self.optimizer.step()

        batch_importance = coords.shape[0] / len(train_data)
        train_epoch_loss += batch_loss.detach().cpu().item() * batch_importance

        # use the self.evaluate to get loss on the validation set
        if (i == (self.epoch - 1)):
            val_epoch_loss = self.evaluate(val_data, draw_plot = True)
        else:
            val_epoch_loss = self.evaluate(val_data)

        # append the losses
        train_loss_list.append(train_epoch_loss)
        val_loss_list.append(val_epoch_loss)

        if early_stop:
            if val_epoch_loss < lowest_val_loss:
                lowest_val_loss = val_epoch_loss
                weights = self.model.state_dict()

    if draw_curve:
        fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
        ax.set_yscale("log")
        # Plot train loss and validation loss
        ax.plot(range(1, len(train_loss_list) + 1), train_loss_list, label='Train Loss')
        ax.plot(range(1, len(val_loss_list) + 1), val_loss_list, label='Val Loss')
        ax.legend()
        ax.set_xlabel("# Batch")
        ax.set_ylabel("Loss")

    if early_stop:
        self.model.load_state_dict(weights)

    return train_loss_list, val_loss_list

def evaluate(self, data, draw_plot=False):

    # init data loader
    data_loader = data.collate(batch_size).cache()

    # init loss function
    loss_func = nn.MSELoss()
    total_loss = 0.0

```

```

if draw_plot:
    true_energies_all = []
    pred_energies_all = []

with torch.no_grad():
    for batch_data in data_loader:

        # compute energies
        species = train_data_batch['species'].to(device)
        coords = train_data_batch['coordinates'].to(device)
        true_energies = train_data_batch['energies'].to(device).float_
        _, pred_energies = model((species, coords))

        # compute loss
        batch_loss = loss_func(true_energies, pred_energies)

        batch_importance = coords.shape[0] / len(data)
        total_loss += batch_loss.detach().cpu().item() * batch_importance

        if draw_plot:
            true_energies_all.append(true_energies.detach().cpu().numpy())
            pred_energies_all.append(pred_energies.detach().cpu().numpy())

if draw_plot:
    true_energies_all = np.concatenate(true_energies_all)
    pred_energies_all = np.concatenate(pred_energies_all)
    # Report the mean absolute error
    # The unit of energies in the dataset is hartree
    # please convert it to kcal/mol when reporting the mean absolute error
    # 1 hartree = 627.5094738898777 kcal/mol
    # MAE = mean(|true - pred|)
    hartree2kcalmol = 627.5094738898777 # changed
    mae = np.mean(np.abs(true_energies_all * hartree2kcalmol - pred_energies_all))
    fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
    ax.scatter(true_energies_all, pred_energies_all, label=f"MAE: {mae}")
    ax.set_xlabel("Ground Truth")
    ax.set_ylabel("Predicted")
    xmin, xmax = ax.get_xlim()
    ymin, ymax = ax.get_ylim()
    vmin, vmax = min(xmin, ymin), max(xmax, ymax)
    ax.set_xlim(vmin, vmax)
    ax.set_ylim(vmin, vmax)
    ax.plot([vmin, vmax], [vmin, vmax], color='red')
    ax.legend()
    #print(mae)

return total_loss

```

```
In [13]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
         model = nn.Sequential(
             aev_computer,
             ani_net
         ).to(device)
```


```
In [14]: model = model


         trainer = ANITrainer(model, learning_rate=1e-3, batch_size = 8192, epoch = 30)
         train_losses, val_losses = trainer.train(train_data, val_data)
```

Sequential - Number of parameters: 197636

Initialize training data...


100% |██████████| 30/30 [04:47<00:00, 9.57s/it]





```
In [15]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[15]: 5.7331837363884256e-05
```

In [16]: `print(train_losses)`

```
[0.004150044536737322, 0.00013599351319533524, 9.514818989270177e-05, 7.027852081559144e-05, 5.559908003337367e-05, 4.661247391919525e-05, 4.0665786512773615e-05, 3.621267203060398e-05, 3.270527230855875e-05, 3.1402471820477994e-05, 2.812861814916419e-05, 2.755122388931103e-05, 2.7531479899193634e-05, 2.9291981469117514e-05, 3.693555873841545e-05, 3.3873674977873154e-05, 5.1292386025111436e-05, 4.623098497962414e-05, 7.731075727434039e-05, 7.103601808456633e-05, 0.00012720442190070373, 2.1829907562654332e-05, 9.596213038330959e-05, 2.6061130264811148e-05, 6.244133583872179e-05, 5.088604436045897e-05, 0.00010542876427378199, 1.8358793291461095e-05, 7.876913400413208e-05, 1.7294839518349943e-05]
```

Checkpoint 3

The lowest loss after checkpoint 2 was 1.7294839518349943e-05, and the MAE was 3.99 kcal/mol for both validation and test data. Now I will optimize the loss and MAE by implementing regularization techniques and performing hyperparameter tuning. My approaches include: modifying the architecture in each atomic net, using more hidden layers, varying the hidden layer size, adding dropouts, changing activation function

Iteration 1

Here I decreased the learning rate from $1e-3$ to $1e-5$

```
In [17]: ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [18]: model = model


# tried decreasing the learning rate
trainer = ANITrainer(model, learning_rate=1e-5, batch_size = 8192, epoch = 30)


train_losses, val_losses = trainer.train(train_data, val_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

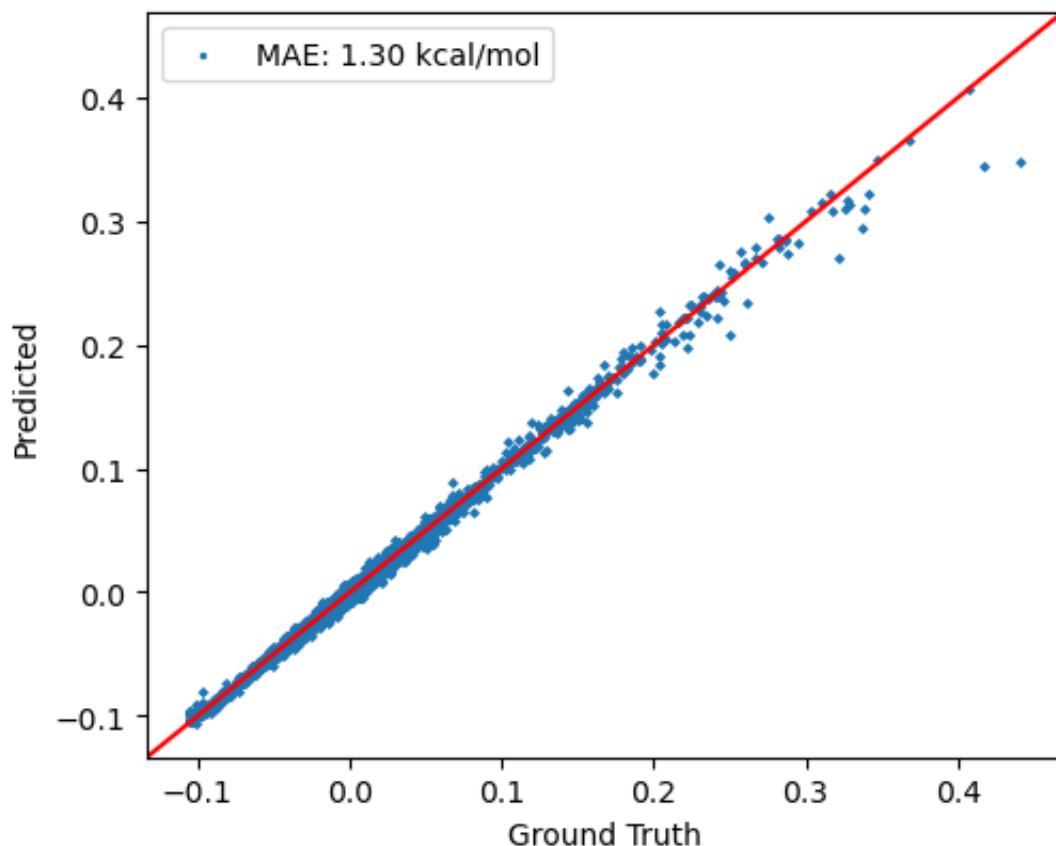
100% |██████████| 30/30 [04:55<00:00, 9.85s/it]





```
In [19]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[19]: 1.1618574713498447e-05
```



```
In [20]: print(train_losses)
```

```
[1.6198787425184687e-05, 1.437100855824389e-05, 1.4284719865594101e-05, 1.4193105494001543e-05, 1.409366437732297e-05, 1.3986305857150458e-05, 1.3877033926014824e-05, 1.3766084715054987e-05, 1.3654304865283119e-05, 1.3544284407360258e-05, 1.3437554973404667e-05, 1.3333548492603435e-05, 1.3228214894807008e-05, 1.3124391283397342e-05, 1.3020091951700905e-05, 1.2920777155121306e-05, 1.2827279716368306e-05, 1.2737543308203968e-05, 1.2647523866215924e-05, 1.2557485454078932e-05, 1.2470214182097147e-05, 1.2385469280158987e-05, 1.2302969178580005e-05, 1.2223159963160752e-05, 1.2145029925637953e-05, 1.2067528187063236e-05, 1.1991593831423151e-05, 1.1917112312281973e-05, 1.1844411974322837e-05, 1.1774181918935158e-05]
```

This gives a very optimal loss of $1.1774181918935158 \times 10^{-5}$, which is even less than what we had achieved in checkpoint 2. Furthermore, we have a validation and test MAE of 1.30 kcal/mol which is significantly less than the MAE from iteration 1 of 3.99. Since the MAE for this iteration is less than 2 kcal/mol, we can see that this model architecture is efficiently and accurately predicts molecular preproperties.

Altering Hidden Layers and Number of Neurons

Iteration 2

Here I change the hidden layer size to be 200 neurons.

```
In [27]: # change hidden layer to 200 neurons
class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 200),
            nn.ReLU(),
            nn.Linear(200, 1)
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()



# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [28]: model = model

trainer = ANITrainer(model, learning_rate=1e-5, batch_size = 8192, epoch = 3


train_losses, val_losses = trainer.train(train_data, val_data)

Sequential - Number of parameters: 308804
Initialize training data...
100%|██████████| 30/30 [05:15<00:00, 10.50s/it]
```



```
In [29]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[29]: 0.0001510852903961897
```



```
In [30]: print(train_losses)
```

```
[0.011198594086009303, 0.004904332405735425, 0.0030942654196581304, 0.0022121939635551344, 0.001655802852550606, 0.0012665693233772789, 0.0009880459464836882, 0.0007894237959531604, 0.0006494236399167477, 0.0005498852287692776, 0.0004771340878089664, 0.0004222732428181937, 0.0003792264057511114, 0.0003441798974062128, 0.00031519307247364125, 0.0002907010556164493, 0.00026971169850084837, 0.00025155130594713403, 0.0002355815441843788, 0.00022139841959921885, 0.000208615520156008, 0.00019700053473018506, 0.00018650730068935957, 0.00017693778144532412, 0.0001681686502599739, 0.00016008687285793452, 0.0001525658371898371, 0.0001456089711597646, 0.0001391919152983602, 0.00013322230234034257]
```

We can see that the lowest loss here is 0.00013322230234034257, which is not as optimal as Iteration 1. Both the validation and test MAE are 4.13.

Iteration 3

Here I add another hidden layer for total of 2 hidden layers

```
In [31]: # change to 2 hidden layers
class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 128),
            nn.ReLU(),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, 1)
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()



# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [32]: model = model

trainer = ANITrainer(model, learning_rate=1e-5, batch_size = 8192, epoch = 3


train_losses, val_losses = trainer.train(train_data, val_data)

Sequential - Number of parameters: 230404
Initialize training data...
100%|██████████| 30/30 [04:56<00:00, 9.88s/it]
```

```
In [33]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[33]: 0.00018131680013287472
```



```
In [34]: print(train_losses)
```

```
[0.016902633900835327, 0.006151463566138277, 0.004080630323489097, 0.0028912382810755764, 0.002164501003997744, 0.0017229005873061597, 0.0014195646698071768, 0.0011846522954481072, 0.0009959653117935375, 0.0008389344439036669, 0.0007092813084363018, 0.0006091732172968584, 0.0005317302496779389, 0.00047127280169100423, 0.0004229998762752826, 0.00038336311315306966, 0.00034985100296456504, 0.0003215336028427507, 0.00029742407860066213, 0.0002764228691111021, 0.00025815981307832217, 0.00024226185530607145, 0.00022838751151681958, 0.00021606824571059532, 0.00020500398995561942, 0.00019495014930457887, 0.00018570129349340076, 0.00017727921520860123, 0.00016954677151862892, 0.00016239312862845572]
```

Similar to iteration 2, iteration 3 leads to a loss of 0.00016239312862845572. This is again not as optimal as iteration 1. The validation and test MAE are 3.91, which is slightly better than the MAE found in iteration 2.

Iteration 4

Now I will change activation function from ReLU to Sigmoid.

```
In [16]: # change activation function
class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 128),
            nn.Sigmoid(),
            nn.Linear(128, 1)
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [17]: model = model



trainer = ANITrainer(model, learning_rate=1e-5, batch_size = 8192, epoch = 30)

train_losses, val_losses = trainer.train(train_data, val_data)
```

Sequential - Number of parameters: 197636


Initialize training data...

100% | ██████████ | 30/30 [05:15<00:00, 10.53s/it]



```
In [18]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[18]: 0.00538318540836897
```



```
In [19]: print(train_losses)
```

```
[1.5109391065437816, 0.5500195810874008, 0.2331469871181487, 0.1459520590001
274, 0.11423158201673402, 0.09482943795281365, 0.0804987498352487, 0.0696012
8407505019, 0.06115456678750528, 0.05439721283410198, 0.048772662073493796,
0.04390613837408042, 0.039564907185506995, 0.035615733834595716, 0.031988203
731551416, 0.02864782305148718, 0.025578042204748432, 0.02276968319424918,
0.020215203202925903, 0.017906103204565784, 0.01583196588893969, 0.013980424
235496773, 0.012337452546964849, 0.010887857790499698, 0.00961575704842344,
0.008504947504170511, 0.0075393399237242194, 0.006703351922989068, 0.0059820
55371318673, 0.005361442608532155]
```

This led to a worse performance than all the previous iterations. The loss is 0.005361442608532155 and the MAE is a very high 28.78 kcal/mol.

Iteration 5

Here I add dropout with the Sigmoid activation function.

```
In [20]: # Add Dropout
class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 128),
            nn.Dropout(p=0.15),
            nn.Sigmoid(),
            nn.Linear(128, 1),
            nn.Dropout(p=0.15),
            nn.Sigmoid(),
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()



# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [21]: model = model

trainer = ANITrainer(model, learning_rate=1e-5, batch_size = 8192, epoch = 3


train_losses, val_losses = trainer.train(train_data, val_data)

Sequential - Number of parameters: 197636
Initialize training data...
100%|██████████| 30/30 [04:55<00:00, 9.85s/it]
```



```
In [22]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[22]: 4.452102335629386
```



```
In [23]: print(train_losses)
```

```
[25.42728441695278, 23.7217000573818, 22.106758819508894, 20.58712591213645
7, 19.163246333882157, 17.8357531841696, 16.600765391986602, 15.457164777947
526, 14.405716801323448, 13.43582456236555, 12.53967674584036, 11.7194286971
87189, 10.966857420888148, 10.273054157884143, 9.63960918208587, 9.054474303
075532, 8.526406179845244, 8.029856606249867, 7.575496110684058, 7.157020076
723713, 6.769853723063323, 6.419699803022335, 6.087143596358469, 5.785909187
165259, 5.504668085305071, 5.244672303925321, 5.001379956716937, 4.778309161
511738, 4.56298850804464, 4.369976060625383]
```

Similar to the last iteration, iteration 5 did not reduce the loss and performed worse than the previous iterations. The loss was 4.369976060625383 and the MAE is the highest it has been, 1244.18 kcal/mol. Evidently, dropout is not an effective regularization technique for this dataset.

Iteration 6

Adding an additional hidden layer and decreasing the batch size and decreasing the learning rate

```
In [24]: # change to 2 hidden layers
class AtomicNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layers = nn.Sequential(
            nn.Linear(384, 128),
            nn.ReLU(),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, 1)
        )

    def forward(self, x):
        return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()



# ANI model requires a network for each atom type
# use torchANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
```

```
In [25]: model = model

trainer = ANITrainer(model, learning_rate=1e-6, batch_size = 512, epoch = 30)


train_losses, val_losses = trainer.train(train_data, val_data)

Sequential - Number of parameters: 230404
Initialize training data...
100%|██████████| 30/30 [05:12<00:00, 10.42s/it]
```



```
In [26]: trainer.evaluate(test_data, draw_plot = True)
```

```
Out[26]: 0.005478806803251806
```



```
In [27]: print(train_losses)
```

```
[0.03929618941770233, 0.034070429020977566, 0.02978522476788199, 0.02630033183955821, 0.02348309789126103, 0.021210439898901142, 0.019370547304545196, 0.017865402058884282, 0.01661317994193476, 0.015548132281081628, 0.014619789697872967, 0.013791077278980724, 0.013036152717752358, 0.012337521710276028, 0.011683877044643815, 0.01106817112568961, 0.010486652314335971, 0.009937363379443571, 0.009418975679104711, 0.008930188727469417, 0.008469817711170233, 0.008036897714337686, 0.007630460965456753, 0.007249356522645049, 0.006892232608928757, 0.006557624200998794, 0.006244065058841623, 0.005950216163739085, 0.005674777854223976, 0.005416390709110152]
```

From iteration 6, we can see that the loss only decreased to 0.005416390709110152 which is still not as optimal as iteration 1. The MAE of 31.18 kcal/mol is again not optimal.

Conclusions

From 6 iterations of modifying the atomic net architecture and tuning the hyperparameters, I can see that what helped optimize the loss most was the learning rate. Decreasing the learning rate helped stabilize oscillation in the loss functions that were seen when the learning rate was $1e-3$. Furthermore, decreasing the learning rate to $1e-5$ helped decrease the overall loss to its lowest point of $1.1774181918935158e-05$ and led to a test/val MAE of 1.30 kcal/mol.

Checkpoint 4

The best model architecture was from Iteration 1, having a learning rate of $1e-5$, batch size of 8192, l2 of $1e-5$ and 30 epochs.

```
In [21]: ani_net = torchani.ANiModel([net_H, net_C, net_N, net_O])
         model = nn.Sequential(
             aev_computer,
             ani_net
         ).to(device)
```

```
In [22]: model = model



         # tried decreasing the learning rate
         trainer = ANITrainer(model, learning_rate=1e-5, batch_size = 8192, epoch = 30)

         train_losses, val_losses = trainer.train(train_data, val_data)
```

Sequential - Number of parameters: 197636

Initialize training data...

100%|██████████| 30/30 [04:47<00:00, 9.58s/it]



In []: